

KESAVA SRAVAN HARIYAPURAJU

Bengaluru, India | +91-9398495822 | kesavasravan467@gmail.com | [GitHub](#) | [LinkedIn](#) | [Portfolio](#)

PROFESSIONAL SUMMARY

Applied AI Specialist with hands-on experience architecting, scaling, and deploying production AI systems and Agentic workflows. Skilled in bridging enterprise customer requirements with robust AI architecture, leveraging LLMs (GPT, Claude, Gemini), Model Context Protocol (MCP), LangGraph, and RAG pipelines powered by FastAPI and cloud containerization. Proven track record of end-to-end technical ownership, enterprise API integrations, vector databases, and production monitoring for high-reliability AI applications.

TECHNICAL SKILLS

AI Engineering: LLMs (GPT-4, Claude, Gemini), Agentic AI, Model Context Protocol (MCP), RAG, LangChain, LangGraph, Prompt Engineering, Semantic Search, FAISS, Hugging Face, Sentence Transformers, NLP, scikit-learn
Languages: Python, Java, SQL, C
Backend & APIs: FastAPI, Spring Boot, REST APIs, Microservices, JWT, RBAC, Node.js
Cloud & DevOps: Docker, Docker Compose, Kubernetes, Jenkins, GitHub Actions, CI/CD, AWS, Azure, GCP, Linux
Databases: PostgreSQL, MySQL, Vector Databases (FAISS)
Observability & Monitoring: Prometheus, Grafana

PROFESSIONAL EXPERIENCE

System Engineer (Forward Deployment & Applied AI)

Feb 2025 – Present

Publicis Sapient

Bengaluru, India

- Architected and deployed production AI agent workflows using **Agentic Studio** and **Model Context Protocol (MCP)** servers, integrating enterprise APIs to automate complex operational workloads.
- Engineered high-performance semantic search pipelines utilizing multi-provider vector embeddings and FAISS vector databases, cutting retrieval latency by **40%** while improving query accuracy.
- Streamlined containerized microservices deployment with Docker, Docker Compose, and Jenkins CI/CD pipelines, driving high availability, zero-downtime releases, and **99.9%** deployment reliability.
- Implemented comprehensive production monitoring and observability infrastructure using **Prometheus** and **Grafana** to track LLM token metrics, API latencies, and agent execution health in real time.
- Partnered directly with cross-functional enterprise stakeholders in a technical consulting capacity to translate business constraints into scalable, production-ready AI solution architectures.

KEY PROJECTS

AI Vector Space Visualization Platform | *React, Vite, Three.js, FastAPI, LangChain, Docker, FAISS*

- Built and deployed a production-ready full-stack AI platform visualizing **100,000+** high-dimensional vector embeddings in interactive 3D using Three.js and WebGL GPU acceleration.
- Implemented dimensionality reduction algorithms (PCA, UMAP, t-SNE) and clustering techniques (KMeans, DBSCAN) for dynamic semantic space analysis and real-time similarity search.
- Integrated multi-provider embedding generation (Google Gemini, OpenAI, Azure OpenAI, Hugging Face via LangChain, and Sentence-Transformers) with automatic hybrid fallback for rate-limit resilience.
- Engineered FAISS-backed similarity search with cosine/Euclidean neighbor visualization, Excel/CSV ingestion, and vector-space export; containerized the full stack with Docker for one-command deployment.

AI Portfolio with Recruiter Chatbot | *React, Vite, FastAPI, LangChain, RAG, Docker, LLMs*

- Designed and deployed an end-to-end AI portfolio site featuring an interactive recruiter chatbot powered by Retrieval-Augmented Generation (RAG) and LangChain [[Live Demo](#)].
- Built custom prompt engineering workflows, conversation memory, and multi-LLM fallback strategies (GPT/Gemini) to guarantee high accuracy, zero hallucination, and real-time response generation.
- Optimized FastAPI REST API endpoints and vector embedding storage to achieve sub-second retrieval latency and **95%+** query resolution accuracy for active recruiters.

EDUCATION

Amrita Vishwa Vidyapeetham

Bachelor of Technology, Electrical & Computer Engineering | CGPA: 7.5/10

Bengaluru, India

Graduated: 2024